

3 Neural Networks: Advanced Architectures

Leon Herrmann

Stefan Kollmannsberger

Chair of Data Engineering in Construction

Bauhaus-Universität Weimar

*Deep Learning in Computational Mechanics – an introductory course,
Herrmann et al. 2025*



website



book



Contents

- 3.1 Fully Connected Neural Network
- 3.2 Forward Propagation
- 3.3 Differentiation
- 3.4 Backpropagation
- 3.5 Activation Function
- 3.6 Learning Algorithm
- 3.7 Regularization of Neural Networks
- 3.8 Approximating the Sine Function
- 3.9.1 Convolutional Neural Networks
- 3.9.2 Graph Neural Networks
- 3.9.6 Recurrent Neural Networks
- 3.9.8 Physics-Inspired Architectures for Dynamics (Hamiltonian & Lagrangian Neural Networks)
- 4 Introduction to Physics-Informed Neural Networks

3.8 Approximating the Sine Function

Goal is to approximate


$$f(x) = \sin(2\pi x), x \in [-1, 1]$$

Neural network architecture

- 2 hidden layers
- 50 neurons each
- Activation function $\sigma(z_i) = \tanh(z_i)$
- $f_{NN} = w_3^T \sigma \left(W_2 \left(\sigma(w_1^T x + b_1) \right) + b_2 \right) + b_3 = \hat{y}$

Data

- Training set contains 40 samples
- Validation set contains 40 samples
- Generated with $y = \sin(2\pi x) + \epsilon, \epsilon = 0.1 \cdot U(-1,1)$
where $U(-1,1)$ is a uniform random distribution in the interval $[-1,1]$
- Test set is the analytical solution



3.8 Approximating the Sine Function

- Goal is to approximate

$$f(x) = \sin(2\pi x), x \in [-1, 1]$$

- Network architecture

$$f_{NN} = w_3^T \sigma \left(W_2 \left(\sigma(w_1^T x + b_1) \right) + b_2 \right) + b_3 = \hat{y}$$

- Cost function used for training

$$C = \mathcal{L}_D^{(\text{train})} = \frac{1}{m_D^{(\text{train})}} \sum_{i=1}^{m_D^{(\text{train})}} \left(\tilde{y}_i^{(\text{train})} - \hat{y}_i^{(\text{train})} \right)^2$$

- Neural network parameter initialization with a uniform distribution with special bounds
- Optimization with Adam for 10'000 epochs

3.8 Approximating the Sine Function – PyTorch

Neural network definition

```
modules = []  
modules.append(torch.nn.Linear(1, 50))  
modules.append(torch.nn.Tanh())  
modules.append(torch.nn.Linear(50, 50))  
modules.append(torch.nn.Tanh())  
modules.append(torch.nn.Linear(50, 1))  
  
model = torch.nn.Sequential(*modules)
```

Training data

```
x = torch.rand((40, 1)) * 2 - 1  
noise = torch.rand(x.shape) * 0.2 - 0.1  
y = torch.sin(2 * torch.pi * x) + noise
```

3.8 Approximating the Sine Function – PyTorch

Cost function definition

```
def costFunction(y, yPred):  
    return torch.mean((y - yPred) ** 2)
```

Prediction and cost function evaluation

```
yPred = model(x)  
cost = costFunction(y, yPred)  
cost.backward()
```

Optimizer definition

```
epochs = 10000  
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
```


3.8 Approximating the Sine Function – PyTorch

Training loop


```
for epoch in range(epochs):  
    optimizer.zero_grad()  
    yPred = model(x)  
    cost = costFunction(y, yPred)  
    cost.backward()  
    optimizer.step()
```

3.8 Approximating the Sine Function – Results

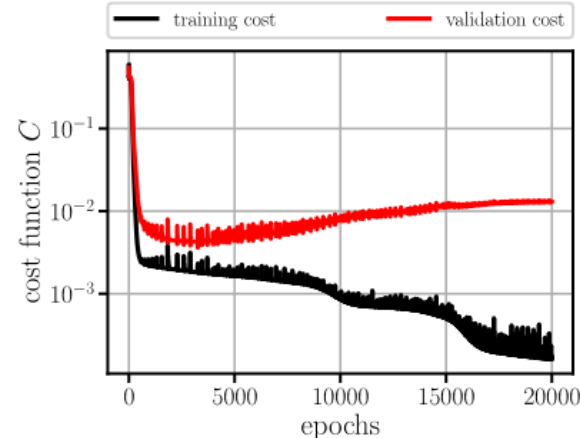

500 epochs



5'000 epochs



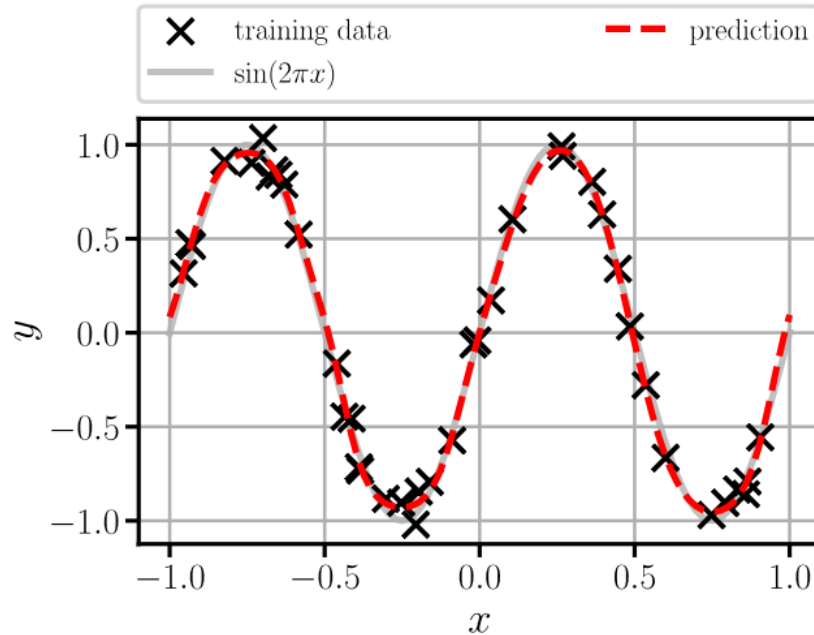
20'000 epochs




Learning history

3.8 Approximating the Sine Function – Results

20'000 epochs with regularization

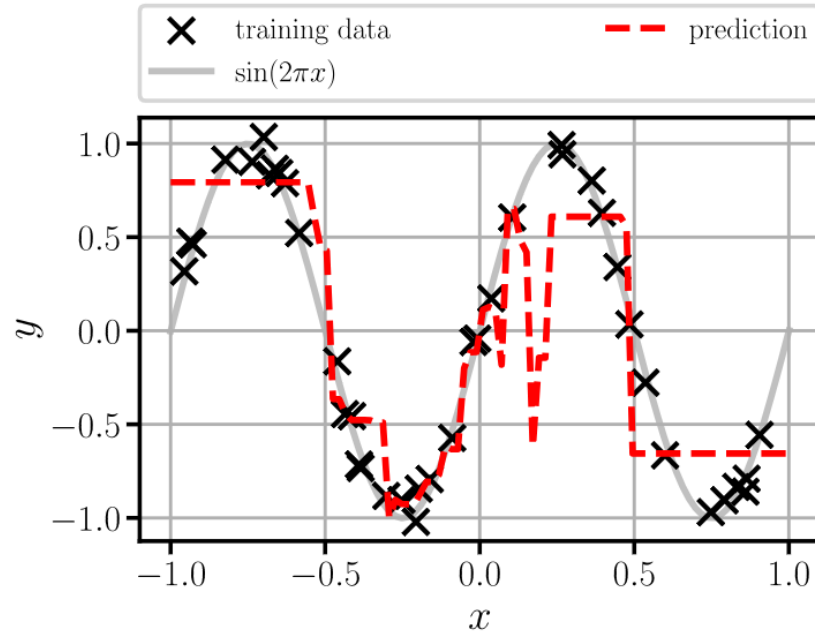


20'000 epochs without regularization

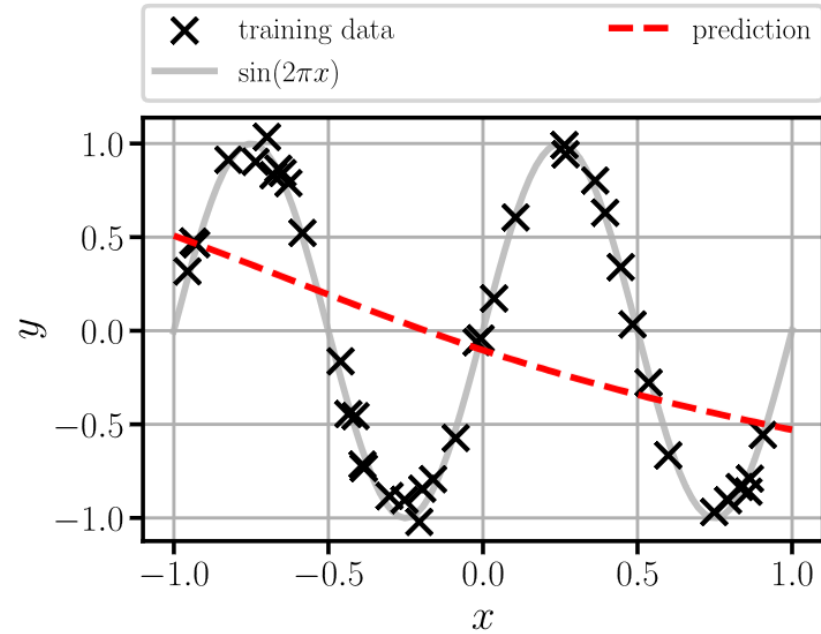


3.8 Approximating the Sine Function – Results

5'000 epochs with Adam at $\alpha = 0.1$



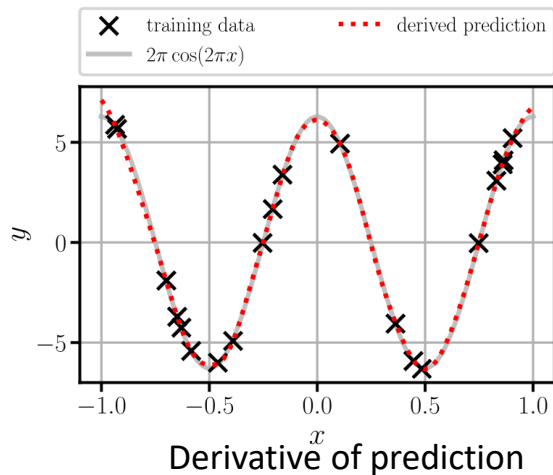


5'000 epochs with SGD



3.8 Approximating the Sine Function – Sobolev

- Derivatives of predictions are typically inaccurate
- Sobolev training** increases accuracy by incorporating derivatives in training data (higher order derivatives also improve)

```
def costFunction(y, yPred, dy, dyPred):
    lossy = torch.mean((y - yPred) ** 2)
    lossdy = torch.mean((dy - dyPred) ** 2)
    return lossy + lossdy
```



Exercises

E.7 Approximating the Sine Function (C)

- In PyTorch, build a fully connected neural network, that approximates the sine function.

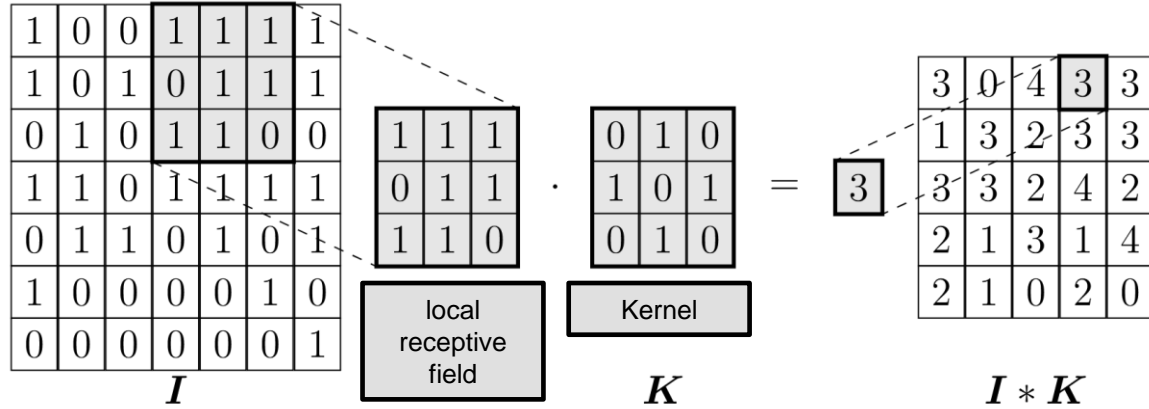
3.9.1 Convolutional Neural Networks

- Originally designed for image processing
- Identify relative positions of features
- Operate with convolutional layers

Convolutional layers

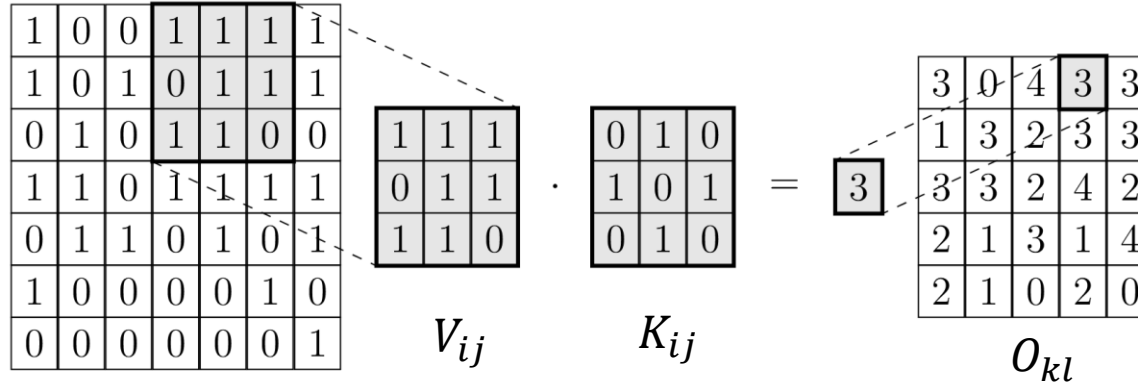
consist of **filters/kernels** $K(w)$ that are applied to an input field I and defined by

- Kernel size
- Number of kernels
- **Stride length**: how large is the shift of the kernel per step
- **Padding**: do you embed the input field?



In neural networks: The values in the kernel are not preset. They are learnt by the network and represent the weights.

3.9.1 Convolutional Neural Networks



The convolution $*$ is computed as a combination of the multidimensional dot product, also known as [tensordot](#) product:

$$o = V_{ij} \cdot K_{ij}$$

Example (2D [convolutional filter](#) applied with a [stride](#) length of $s = 1$ and no padding)

$$1 \cdot 0 + 1 \cdot 1 + 1 \cdot 0 + 0 \cdot 1 + 1 \cdot 0 + 1 \cdot 1 + 1 \cdot 0 + 1 \cdot 1 + 0 \cdot 0 = 3$$

This operation is then repeated by shifting the [view](#) V_{ij} with the stride s

$$O_{kl} = \sum_{i,j} V_{(i+k)(j+l)} \cdot K_{ij} = V * K$$

3.9.1 Convolutional Neural Networks – Filters

Example filters

Identity filter




$$K = \begin{pmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{pmatrix}, s=1$$

Gaussian filter



$$K = \frac{1}{16} \begin{pmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{pmatrix}, s=1$$

Sobel filter in x -direction




$$K = \begin{pmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{pmatrix}, s=1$$

Averaging filter



$$K = \frac{1}{9} \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}, s=3$$

Gaussian Distribution




For other filters for gradient computation, see Chapter 9.

finite difference approximation	convolutional filter
$f'(x) = \frac{f(x+\Delta x) - f(x)}{\Delta x}$	$\frac{1}{\Delta x} [-1, 1]$
$f''(x) = \frac{f(x+\Delta x) - 2f(x) + f(x-\Delta x)}{\Delta x^2}$	$\frac{1}{\Delta x^2} [1, -2, 1]$

3.9.1 Convolutional Neural Networks – Pooling

Pooling layers reduce the spatial size by extracting relevant features (filtering!)

- **Pooling types**
 - Max Pooling
 - Average Pooling
- **Pooling layers are defined by**
 - Kernel size
 - Stride length
 - Padding




- Example “max pooling”(2D convolutional filter applied with a stride length of $s = 1$ and no padding)
 $\max\{1, 2, 1, 2\} = 2$

3.9.1 Convolutional Neural Networks – Example

What if we want two features O from one input?

Example with 1D convolutional filters

- Stride length $s = 1$
- Without a bias
- Input (given)



$$I = [0, 1, 3, 2, 5]$$

- Weights from input layer I to first hidden layer $H^{(1)}$ (given)

$$w_{11}^{(1)} = [0, 2, 1], w_{12}^{(1)} = [1, 0, 3]$$

- Forward pass to first hidden layer $H_j^{(1)} = \sum_i I_i * w_{ij}^{(1)}$


$$H_1^{(1)} = [0 \cdot 0 + 1 \cdot 2 + 3 \cdot 1, 1 \cdot 0 + 3 \cdot 2 + 2 \cdot 1, 3 \cdot 0 + 2 \cdot 2 + 5 \cdot 1] = [5, 8, 9]$$

$$H_2^{(1)} = [0 \cdot 1 + 1 \cdot 0 + 3 \cdot 3, 1 \cdot 1 + 3 \cdot 0 + 2 \cdot 3, 3 \cdot 1 + 2 \cdot 0 + 5 \cdot 3] = [9, 7, 18]$$

3.9.1 Convolutional Neural Networks – Example

Example with 1D convolutional filters

- Stride length $s = 1$
- Without a bias
- Hidden layer $\mathbf{H}^{(1)}$ (given)



$$\mathbf{H}_1^{(1)} = [5, 8, 9]$$

$$\mathbf{H}_2^{(1)} = [9, 7, 18]$$

- Weights from first hidden layer $\mathbf{H}^{(1)}$ to output layer $\mathbf{O}_k = \sum_j \mathbf{H}_j^{(1)} * \mathbf{w}_{jk}^{(2)}$ (given)

$$\mathbf{w}_{11}^{(2)} = [1, 0, 0], \mathbf{w}_{21}^{(2)} = [2, 1, 0]$$

$$\mathbf{w}_{12}^{(2)} = [0, 0, 3], \mathbf{w}_{22}^{(2)} = [0, 1, 1]$$

- Forward pass to output layer

$$\mathbf{O}_1 = [(5 \cdot 1 + 8 \cdot 0 + 9 \cdot 0) + (9 \cdot 2 + 7 \cdot 1 + 18 \cdot 0)] = 30$$

$$\mathbf{O}_2 = [(5 \cdot 0 + 8 \cdot 0 + 9 \cdot 3) + (9 \cdot 0 + 7 \cdot 1 + 18 \cdot 1)] = 52$$

3.9.1 Convolutional Neural Networks – Depth

ImageNet Classification with Deep Convolutional Neural Networks, (Alex) Krizhevsky et al. 2012

Comparison to fully connected neural network

- Channels are analogous to neurons
- Filters are analogous to linear transformations (connections)
- Number of parameters in fully connected neural network:
- Number of parameters in convolutional neural network:


Number of parameters per weight/kernel

Number of connections (weights)

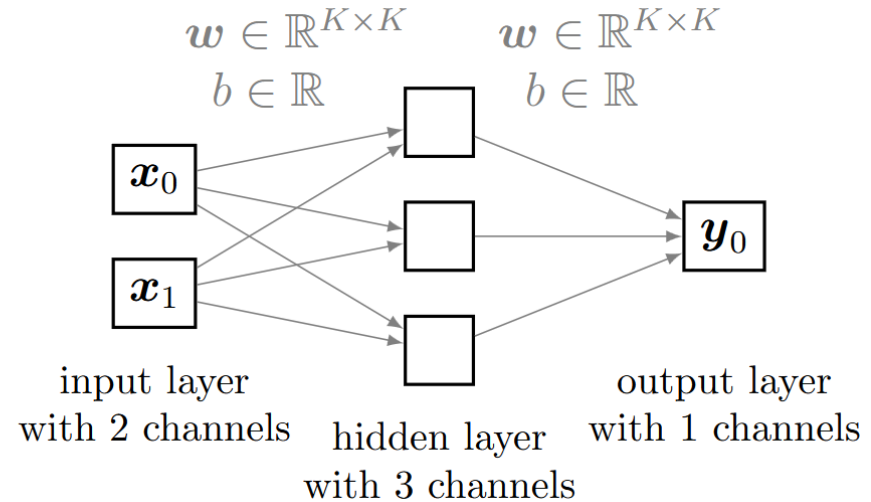
Number of biases

$$1 \cdot (2 \cdot 3 + 3 \cdot 1) + (3 + 1)$$

$$(K \cdot K) \cdot (2 \cdot 3 + 3 \cdot 1) + (3 + 1)$$



Fully connected neural network



Convolutional neural network


3.9.1 Convolutional Neural Networks – Padding

Padding counteracts the shrinking of the output, padding can be used

- Padding types
 - Valid padding (No padding)
 - Zero padding
 - Reflection padding
- The output dimension O after a convolutional layer

$$O = \left\lfloor \frac{I - K + 2P}{s} \right\rfloor + 1$$

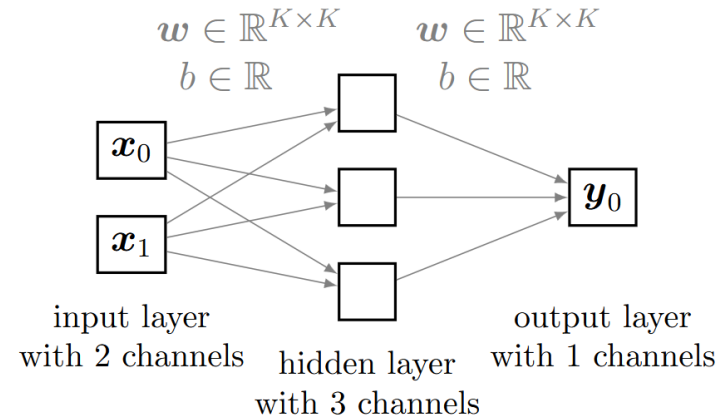
- Input dimension I
- Kernel size K
- Padding size P
- Stride length s



3.9.1 Convolutional Neural Networks – PyTorch

```
class CNN(torch.nn.Module):  
    def __init__(self, inputImages, hiddenImages, outputImages):  
        super(CNN, self).__init__()  
        self.cnn1 = torch.nn.Conv2d(inputImages, hiddenImages, kernel_size=3,  
                                     stride=1, padding=1)  
        self.cnn2 = torch.nn.Conv2d(hiddenImages, outputImages, kernel_size=3,  
                                     stride=1, padding=1)  
        self.activation = torch.nn.ReLU()  
  
    def forward(self, x):  
        y = self.activation(self.cnn1(x))  
        y = self.cnn2(y)  
        return y
```

```
model = CNN(2, 3, 1) # input has size  
                  (1, 2, imageSize, imageSize)
```



Exercises

E.11 Convolutional Neural Network (P & C)

- Familiarize yourself with convolutional neural networks through pen-and-paper forward propagation and through the application of different convolutional filters to images (using PyTorch).



3.9.1 Convolutional Neural Networks – Example

Learning strain distributions from data

- Goal is to make the neural network predict the strain of a 2D domain under a uni-axial load
- Domain is defined by fiber material (red) and matrix material (white) with the elastic properties

$$E_f = 85'000, \nu_f = 0.22$$

$$E_m = 3'000, \nu_m = 0.4$$



3.9.1 Convolutional Neural Networks – Example


- To reduce the complexity of the task, the domain is discretized in 32×32 grid
- Strain prediction from the underlying material distribution
- As neural network architecture, a U-Net (see Chapter 8) is used

U-Net: Convolutional Networks for Biomedical Image Segmentation, Ronneberger et al. 2015





- When training a surrogate model, the **breakeven threshold** τ must be considered

$$\tau = \frac{T_{\text{data}} + T_{\text{train}}}{T_{\text{simulation}} + T_{\text{surrogate}}}$$

- Which represents the number of surrogate evaluations to make training worth it
- T is the time of
 - Data collection (T_{data})
 - Training (T_{train})
 - 1 Simulation ($T_{\text{simulation}}$)
 - 1 Surrogate evaluation ($T_{\text{surrogate}}$)





Assuming $T_{\text{data}} = T_{\text{simulation}} \cdot m$
and $T_{\text{train}} \approx 0$




3.9.1 Convolutional Neural Networks – Example

Model selection (make sure that model can learn the training data by heart)

- NN 1: U-net with 5 convolutional layers: 94'978 parameters
- NN 2: feed-forward convolutional neural network with 4 layers: 38'405 parameters
- NN 3: U-net followed by 3 convolutional layers: 154'471 parameters
- Training with 1 datapoint





ϵ_{11}
Prediction with NN 2



ϵ_{11}
Ground truth



3.9.1 Convolutional Neural Networks – Example

With selected model (NN 2), increase the training data (and potentially tune the architecture & optimizer)



3.9.1 Convolutional Neural Networks – Example

Prediction on unseen data (i.e., validation data)



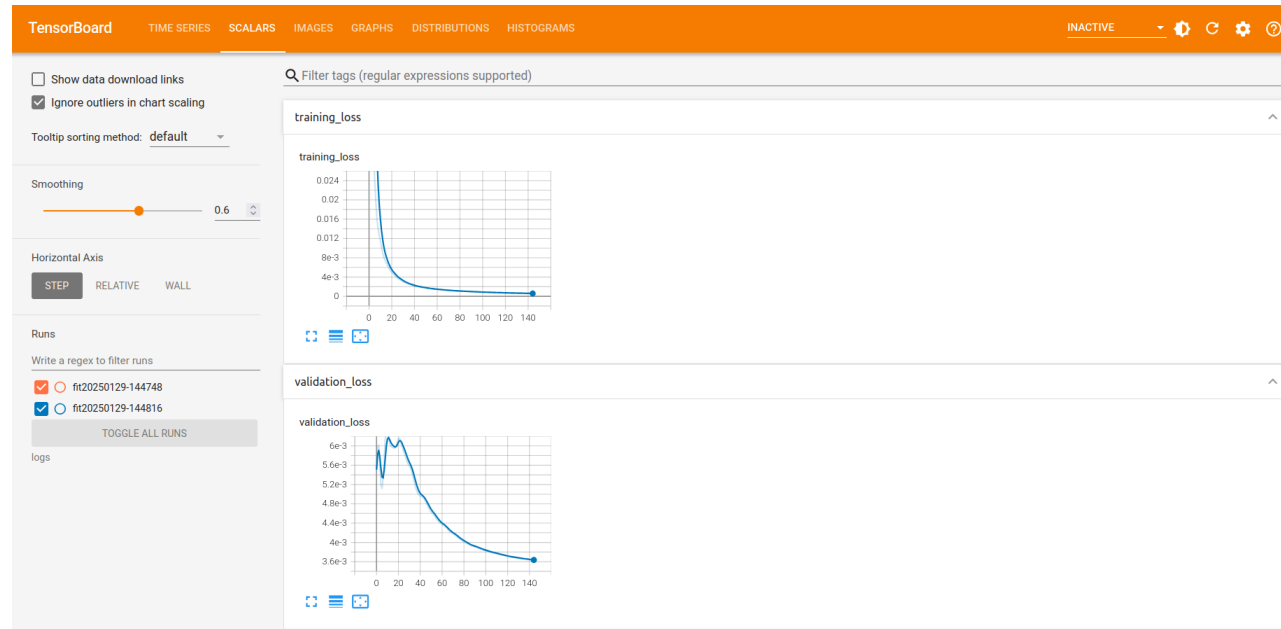
3.9.1 Convolutional Neural Networks – TensorBoard

TensorBoard is a **visualization tool** for TensorFlow (& other machine learning libraries, e.g., PyTorch)

- Beneficial for **debugging** and **hyperparameter tuning/model selection**
- Includes features such as **graphs**, **histograms**, **output visualization**

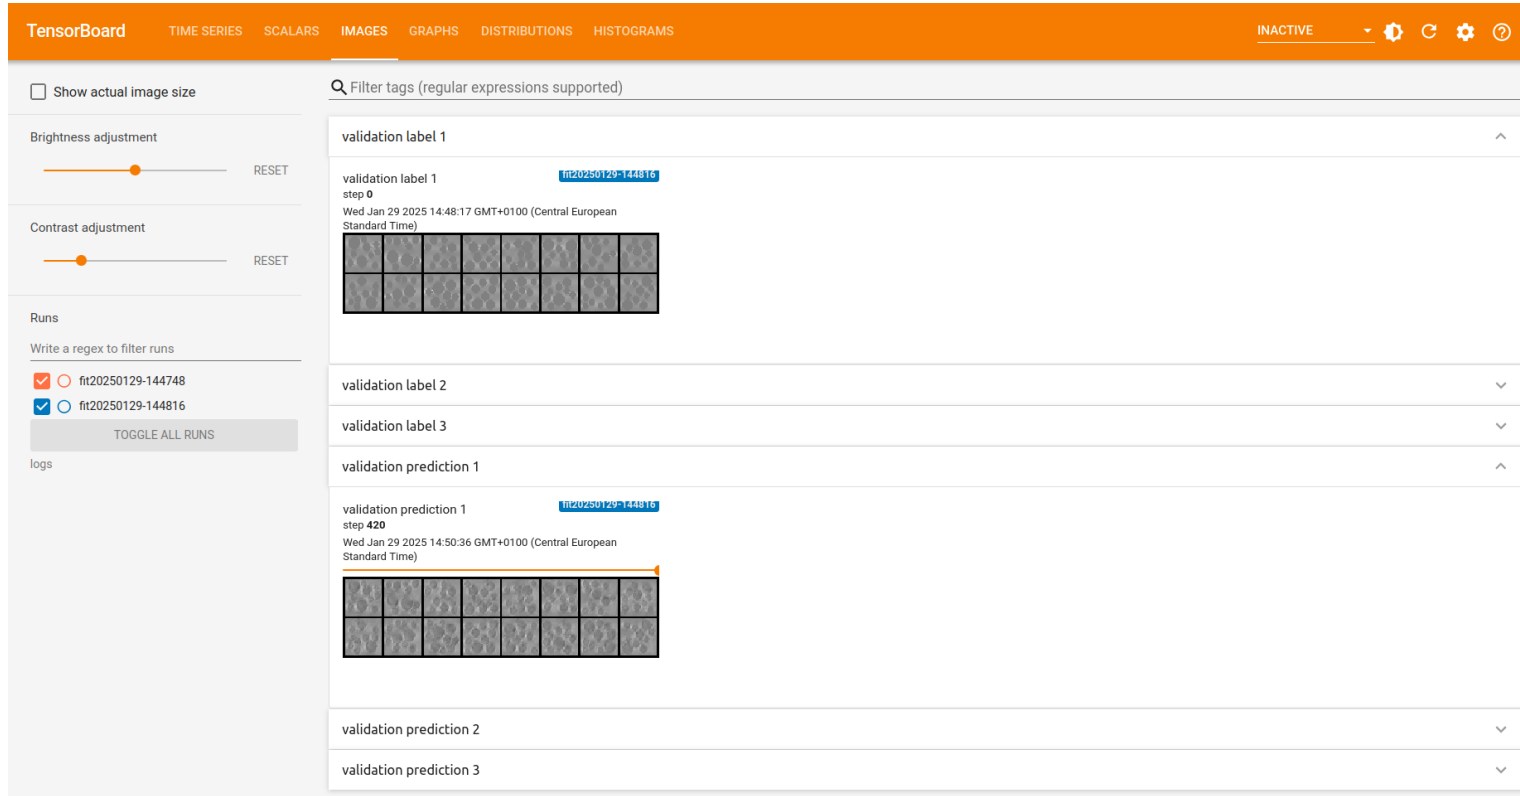


TensorBoard




3.9.1 Convolutional Neural Networks – TensorBoard

Track the progress in prediction quality by comparing outputs with labels



3.9.1 Convolutional Neural Networks – TensorBoard

Compare different architectures and hyperparameters through different training runs



Exercises


E.12 Learning Strain Distributions (C)

- Train a convolutional neural network as a surrogate, that maps the Young's modulus distribution to strain distributions. Experiment with different hyperparameters and different network architectures. For this, use TensorBoard.

Normalization

Normalization helps the optimization procedure

- Ensures, that features have similar scales



- Input normalization centering the mean to zero and establishing a unit standard deviation

$$\hat{x}_i = \frac{x_i - \mu}{\sigma}$$

- where

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i, \sigma = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2}$$

Batch Normalization

Batch Normalization normalizes the activations between the layers to help the optimization

$$y = \frac{x - \mu(x)}{\sigma(x) + \epsilon} \cdot \gamma + \beta$$

Batch Normalization is composed of two parts

- a) Centering the mean around zero and normalizing the standard deviation to a unit standard deviation

$$\tilde{x} = \frac{x - \mu(x)}{\sigma(x) + \epsilon}$$

- b) scaling and shifting by an element-wise modification with the trainable parameters γ, β

$$y = \tilde{x} \cdot \gamma + \beta$$

- This increases the model capacity further and makes it possible for the network to learn an optimal normalization

Batch normalization is:

- typically applied after the activation function, but before is also possible
- Almost always encountered when using convolutional neural networks (alternate normalizations possible)


In PyTorch it can be defined with `torch.nn.BatchNorm1d` OR `torch.nn.BatchNorm2d`

3.9.2 Graph Neural Networks

- Convolutional neural networks are limited to data aligned on rectangular & uniform grids
- Graph neural networks** are a generalization to general graphs

- Graph**

- Nodes/vertices v_i
- Edges e_i
- Global attribute u



- Many different architectures exist, ([message passing networks](#), [graph convolutional networks](#), [graph transformers](#)...)
- We will focus on **message passing networks**, which consider the following invariant

$$v_i^s \xrightarrow{e_i} v_i^r$$

Relational inductive biases, deep learning, and graph networks, Battaglia et al. 2018

- Each edge e_i has a sender v_i^s and receiver node v_i^r

3.9.2 Graph Neural Networks

- Edges are updated through the **first** fully connected neural network $e'_i = f^e(e_i, v_i^s, v_i^r, u)$
- Nodes are updated by **aggregating** connecting edges \bar{e}_i (e.g., max or sum): **second** fully connected neural network $v'_i = f^v(\bar{e}', v_i, u)$
- The global attribute is updated by **aggregating** all edges and nodes: **third** fully connected neural network $u' = f^u(\bar{e}', \bar{v}', u)$

Algorithm 8 Graph block in a message passing neural network

Require: Graph consisting of nodes \mathbf{v} , edges \mathbf{e} and a global attribute u .

for all edges e_i **do**

 Update edges: $e'_i = f^e(e_i, v_i^s, v_i^r, u; \Theta_e)$

end for

for all nodes v_i **do**

 Find all edges connecting to node v_i : e'_i

 Aggregate adjacent edges: $\bar{e}'_i = \rho^{e \rightarrow v}(e'_i)$

 Update nodes: $v'_i = f^v(\bar{e}'_i, v_i, u; \Theta_v)$

end for

Aggregate all edges \mathbf{e} : $\bar{e}' = \rho^{e \rightarrow u}(\mathbf{e})$

Aggregate all nodes \mathbf{v} : $\bar{v}' = \rho^{v \rightarrow u}(\mathbf{v})$

Update global attribute: $u' = f^u(\bar{e}', \bar{v}', u; \Theta_u)$


- Application to mesh-based simulation

Learning Mesh-Based Simulation with Graph Networks, Pfaff et al. 2020

3.9.6 Recurrent Neural Networks

Recurrent Neural Networks reuse information from previous states by

- looping over itself to generate sequences.



- Equivalent to using multiple copies of the same network with a **hidden state** h_t (unrolled recurrent neural network)

$$h_t = \sigma(W_x x_t + b_x + W_h h_{t-1} + b_h)$$

- W_x, W_h are the weights and b_x, b_h the biases
- y_t is obtained by a (learnable) mapping between h_t and y_t
- Useful for applications with **sequential data**, such as speech recognition, language modeling, translation
- In practice, only able to connect recent information with each other, e.g., "*We used too many lemons for our lemonade ... We did not like the lemonade, it was too sour*" poses a problem, as the cause for the is too far from the outcome
- Overcome by long short-term memory networks (LSTMs), gated recurrent units (GRUs), and transformers

3.9.6 Recurrent Neural Networks - PyTorch

```
class RNN(torch.nn.Module):  
    def __init__(self, inputSize, hiddenStateSize, outputSize):  
        super(RNN, self).__init__()  
        self.rnn = torch.nn.RNN(inputSize, hiddenStateSize, nonlinearity='relu',  
num_layers=2)  
        self.linear = torch.nn.Linear(hiddenStateSize, outputSize)  
  
    def forward(self, x):  
        h, _ = self.rnn(x) # unrolls the RNN  
        y = self.linear(h) # transforms hidden layer to output layer  
        return y
```

```
inputSize = 1  
hiddenStateSize = 500  
outputSize = 1  
model = CNN(inputSize, hiddenStateSize, outputSize)
```


Recurrent neural networks do not only have to be combined with fully connected neural networks, convolutional neural networks are for example also possible

3.9.8 Physics-Inspired Architectures for Dynamics

- The previous architecture learn physical behavior only from data
- Governing laws can be incorporated into neural networks
 - Through training (via penalty terms in the cost function, i.e., [weak enforcement](#)/regularization; see Chapters 4 & 5)
 - Via the neural network architecture (by constraining the learnable space, i.e., [strong enforcement](#); see Chapter 7)
- Strong enforcement of physics
 - Dynamics
 - **Hamiltonian neural networks**
 - **Lagrangian neural networks**
 - Constitutive modeling
 - Input-convex neural networks (see Chapter 7)
 - Boundary conditions
 - Strong enforcement of boundary conditions (see Chapters 4 & 5)

3.9.8 Physics-Inspired Architectures for Dynamics

Consider a one-dimensional mass-spring model



Described by the [ordinary differential equation](#)

$$m\ddot{x}(t) + kx(t) = 0$$

And the solution

$$x(t) = A\sin(\omega t + \phi)$$

Where $\omega = \sqrt{k/m}$ is the natural frequency and A, ϕ are determined by the initial conditions $x(0), \dot{x}(0)$

Alternatively the system can be described by

- [Hamiltonian Mechanics](#)
- [Euler-Lagrange Equation](#)

3.9.8.2 Hamiltonian Mechanics

In Hamiltonian mechanics, systems are described by **coordinate pairs** (position & momentum, i.e., $\mathbf{p}(t) = m\dot{\mathbf{x}}(t)$)
 $[\mathbf{x}(t), \mathbf{p}(t)]$

Scalar **Hamiltonian** $\mathcal{H}(\mathbf{x}(t), \mathbf{p}(t))$ represents the system's total energy Π_{tot} and fulfills

$$\begin{aligned}\frac{d\mathbf{x}}{dt} &= \frac{\partial \mathcal{H}}{\partial \mathbf{p}}, \\ \frac{d\mathbf{p}}{dt} &= -\frac{\partial \mathcal{H}}{\partial \mathbf{x}}\end{aligned}$$

Given the Hamiltonian, the time evolution of \mathbf{x}, \mathbf{p} can be computed via **integration**

$$\begin{aligned}\mathbf{x}(t + \Delta t) &= \mathbf{x}(t) + \int_t^{t+\Delta t} \left. \frac{\partial \mathcal{H}}{\partial \mathbf{p}} \right|_{\tau} d\tau \\ \mathbf{p}(t + \Delta t) &= \mathbf{p}(t) - \int_t^{t+\Delta t} \left. \frac{\partial \mathcal{H}}{\partial \mathbf{x}} \right|_{\tau} d\tau\end{aligned}$$

Which can, e.g., be discretized with a **forward Euler** scheme

$$\mathbf{x}(t + \Delta t) = \mathbf{x}(t) + \left. \frac{\partial \mathcal{H}}{\partial \mathbf{p}} \right|_t \Delta t, \quad \mathbf{p}(t + \Delta t) = \mathbf{p}(t) - \left. \frac{\partial \mathcal{H}}{\partial \mathbf{x}} \right|_t \Delta t$$

3.9.8.3 Hamiltonian Neural Networks

For the one-dimensional mass-spring model, the **Hamiltonian** is known as

$$\mathcal{H} = \Pi_{\text{tot}} = \Pi_{\text{kin}} + \Pi_{\text{pot}} = \frac{1}{2}m\dot{x}^2 + \frac{1}{2}kx^2$$

But what if the Hamiltonian is unknown, but the system must obey Hamiltonian mechanics?


Hamiltonian neural networks

- **Data triplets** $\{\tilde{\mathbf{x}}(t_i), \tilde{\dot{\mathbf{x}}}(t_i), \tilde{\mathbf{p}}(t_i)\}_{i=1}^m$ (enable computation of $\tilde{\mathbf{p}}(t_i), \tilde{\dot{\mathbf{p}}}(t_i)$)
- Mapping to the **Hamiltonian** is to be learned $\hat{\mathcal{H}} = \mathcal{H}(\mathbf{x}(t), \mathbf{p}(t); \Theta)$
- From $\hat{\mathcal{H}}$, $d\hat{\mathbf{x}}/dt$ and $d\hat{\mathbf{p}}/dt$ can be computed $\frac{d\mathbf{x}}{dt} = \frac{\partial \mathcal{H}}{\partial \mathbf{p}}, \frac{d\mathbf{p}}{dt} = -\frac{\partial \mathcal{H}}{\partial \mathbf{x}}$
- A comparison of $d\hat{\mathbf{x}}/dt$ and $d\hat{\mathbf{p}}/dt$ to the corresponding data in the data triplets, enables a **cost function**

$$\mathcal{C} = \frac{1}{m} \sum_{i=1}^m \left(\left\| \frac{d\tilde{\mathbf{x}}(t_i)}{dt} - \frac{\partial \hat{\mathcal{H}}(\tilde{\mathbf{x}}(t_i), \tilde{\mathbf{p}}(t_i); \Theta)}{\partial \tilde{\mathbf{p}}(t_i)} \right\|^2 + \left\| \frac{d\tilde{\mathbf{p}}(t_i)}{dt} + \frac{\partial \hat{\mathcal{H}}(\tilde{\mathbf{x}}(t_i), \tilde{\mathbf{p}}(t_i); \Theta)}{\partial \tilde{\mathbf{x}}(t_i)} \right\|^2 \right)$$

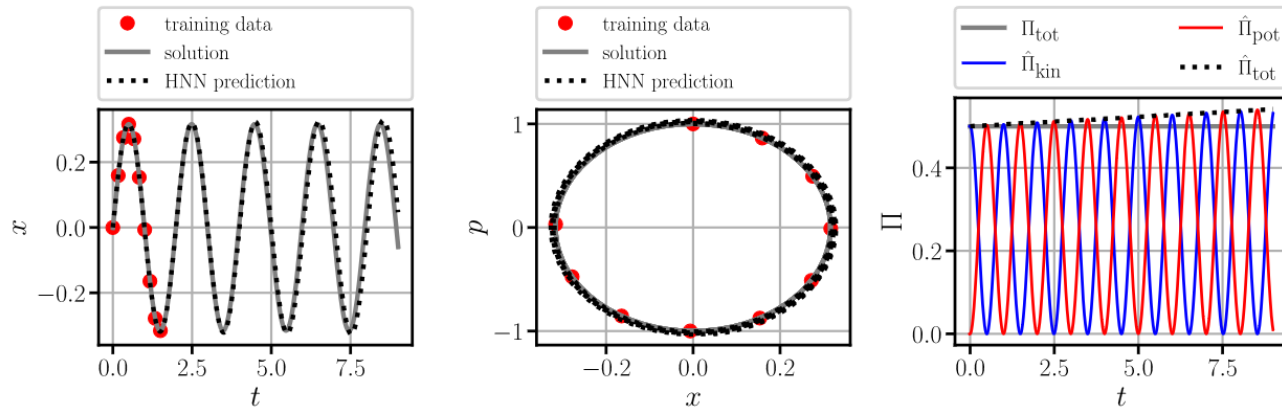
- Prediction of the trajectory $\hat{\mathbf{x}}(t)$ beyond the datapoints seen in the data triplets, possible through **forward Euler**

$$\mathbf{x}(t + \Delta t) = \mathbf{x}(t) + \frac{\partial \mathcal{H}}{\partial \mathbf{p}} \bigg|_t \Delta t, \quad \mathbf{p}(t + \Delta t) = \mathbf{p}(t) - \frac{\partial \mathcal{H}}{\partial \mathbf{x}} \bigg|_t \Delta t$$



3.9.8.3 Hamiltonian Neural Networks - Results


Hamiltonian neural network



Standard neural network

$$\begin{pmatrix} \hat{\tilde{x}} \\ \hat{\tilde{p}} \end{pmatrix} = f(x, p; \Theta)$$

$$C = \frac{1}{m} \sum_{i=1}^m (\tilde{x} - \hat{\tilde{x}})^2 + (\tilde{p} - \hat{\tilde{p}})^2$$



3.9.8.4 Euler-Lagrange Equation

The **Lagrangian framework** offers a more general framework than Hamiltonian mechanics

The **Lagrangian** of the mass-spring system is defined as

$$L = \Pi_{\text{kin}} - \Pi_{\text{pot}} = \frac{1}{2} m \dot{x}^2 - \frac{1}{2} k x^2$$


Which obeys the **Euler-Lagrange equation** (ensuring the principle of least action)

$$\frac{d}{dt} \frac{\partial L}{\partial \dot{x}} = \frac{\partial L}{\partial x}$$

The acceleration \ddot{x} can be obtained by rewriting the Euler-Lagrange equation

$$\begin{aligned} \frac{d}{dt} \frac{\partial L}{\partial \dot{x}} &= \frac{\partial^2 L}{\partial x \partial \dot{x}} \dot{x} + \frac{\partial^2 L}{\partial \dot{x}^2} \ddot{x} = \frac{\partial L}{\partial x} \\ \ddot{x} &= \left(\frac{\partial^2 L}{\partial \dot{x}^2} \right)^{-1} \left(\frac{\partial L}{\partial x} - \frac{\partial^2 L}{\partial x \partial \dot{x}} \dot{x} \right) \end{aligned}$$


Similar as for the Hamiltonian mechanics, the evolution of $x(t)$ can be computed by **integrating** \dot{x}, \ddot{x}



3.9.8.5 Lagrangian Neural Networks





- **Data triplets** $\{\tilde{x}(t_i), \tilde{\dot{x}}(t_i), \tilde{\ddot{x}}(t_i)\}_{i=1}^m$
- Mapping to the **Lagrangian** is to be learned $\hat{L} = L(x(t), \dot{x}(t); \Theta)$
- The **acceleration** $\hat{\ddot{x}}$ can be computed from
$$\ddot{x} = \left(\frac{\partial^2 L}{\partial \dot{x}^2} \right)^{-1} \left(\frac{\partial L}{\partial x} - \frac{\partial^2 L}{\partial x \partial \dot{x}} \dot{x} \right)$$
- With the predicted acceleration a **cost function** can be formulated as

$$C = \frac{1}{m} \sum_{i=1}^m \left\| \tilde{\ddot{x}}(t_i) - \hat{\ddot{x}} \left(\hat{L}; \tilde{x}(t_i), \tilde{\dot{x}}(t_i) \right) \right\|^2$$



3.9.8.5 Lagrangian Neural Networks - Results

Lagrangian neural network






The Lagrangian is not unique

Standard neural network

$$\begin{pmatrix} \hat{\tilde{x}} \\ \hat{\tilde{p}} \end{pmatrix} = f(x, p; \Theta)$$

$$C = \frac{1}{m} \sum_{i=1}^m (\tilde{x} - \hat{\tilde{x}})^2 + (\tilde{p} - \hat{\tilde{p}})^2$$



Exercises

E.13 Hamiltonian & Lagrangian Neural Networks (C)

- Implement a Hamiltonian and a Lagrangian neural network and apply it to the mass-spring model. Compare the generalization capabilities to a standard neural network.

Contents

- 3.1 Fully Connected Neural Network
- 3.2 Forward Propagation
- 3.3 Differentiation
- 3.4 Backpropagation
- 3.5 Activation Function
- 3.6 Learning Algorithm
- 3.7 Regularization of Neural Networks
- 3.8 Approximating the Sine Function
- 3.9.1 Convolutional Neural Networks
- 3.9.2 Graph Neural Networks
- 3.9.6 Recurrent Neural Networks
- 3.9.8 Physics-Inspired Architectures for Dynamics (Hamiltonian & Lagrangian Neural Networks)
- 4 Introduction to Physics-Informed Neural Networks

3 Neural Networks: Advanced Architectures

Leon Herrmann

Stefan Kollmannsberger

Chair of Data Engineering in Construction

Bauhaus-Universität Weimar

*Deep Learning in Computational Mechanics – an introductory course,
Herrmann et al. 2025*



website



book

